



- Goals and topics

- Goals

- introduce Docker
- show you how to run a Spring Boot application in Docker

- Subjects

- Docker
- Docker in practice
- how Docker works
- a containerized application
- discussion



* Docker



- **Docker** (www.docker.com) is a container platform to build, release and run distributed applications - in a simple, fast, scalable and portable way

- a **Docker container** is a standardized software unit, which packages a software service, along with its configurations and dependencies
 - a container contains everything needed to run that software service - executable code, configurations, libraries, and system tools
- Docker containers are lightweight (they use few resources and boot quickly), standardized and open (and therefore portable: they can run with major Linux distributions and with Windows and Mac OS, and even in the cloud) and secure




History

- Docker platform (2013) was initially built on top of LXC containers (2008)
 - LXC offers a set of kernel features for container management - which are low-level and often difficult to use directly.
- Docker built on this foundation to provide a more powerful and easier-to-use set of high-level tools and features
- Docker today is based on libraries *containerd* And *runc* (2014, 2015) - as well as on *cgroup* And *namespace*
- Docker was an instant hit and is used in production by many companies - few technologies have seen such an adoption rate



Docker

- The platform *Docker* allows for a separation between applications and execution infrastructure
 - to simplify the release of applications
 - to ensure the portability of services implemented through containers - both on premises and in the cloud






Docker Engine

- The fundamental core of the Docker platform is *Docker Engine*

the client can be a GUI or a CLI



Docker Engine

- Docker Engine is based on a client-server architecture
 - the *server* is a host capable of running and managing Docker containers
 - runs the Docker daemon process (*dockerd*)
 - manages a set of Docker objects - containers, images, networks and volumes
 - the *client* (*docker*) accepts commands from the user via a CLI interface and communicates with the Docker daemon on the host
 - communication takes place via a REST API
 - the *registry* contains a set of images
 - Docker's public registry is Docker Hub



Containers and images

- Two basic types of Docker objects
 - a *container* it is, in fact, a container instance, which contains an application or a service - along with everything needed to run it
 - it's a concept *dynamic*, runtime
 - can be run on a host
 - a *image* is a template for creating containers
 - it's a concept *static*
 - it cannot be done directly
 - relationship between container and images
 - each container is created from an image
 - from one image it is possible to create many containers



Images

- In practice, a '*image*' is a set of files - representing a container's file system snapshot
 - eg, an image with an Ubuntu OS, Open JDK, and a specific Java application of interest
 - another image could be specific to NGINX or to Apache Kafka
- an image is a concept *static*, inert
 - it is not done directly
 - it has no state of its own
 - it is immutable



Container

- A **container** is an executable instance of container, created from a Docker image
 - an "application container" - which contains an application or service
 - for example, a distributed software system might include
 - N containers that are all replicas of a web application of interest (based on the same image)
 - an additional container to distribute client requests among the N replicas of the web application of interest (based on an image for NGINX)
- a container is a concept **dynamic**, runtime
 - can be run on a host
 - has its own state - which can change during execution
 - e.g., the contents of the file system (in the disk) or the state of the sessions (in the main memory)

11

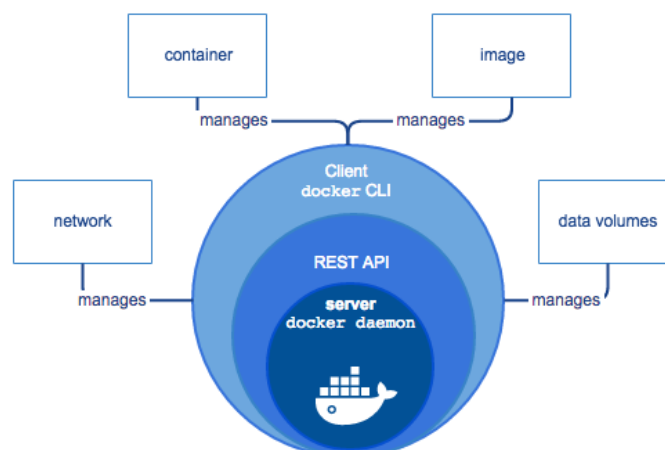
Docker

Luca Cabibbo ASW



The Docker server

- To summarize, the Docker server
 - runs the Docker daemon process
 - manages a set of Docker objects - mostly containers and images
- allows access to its clients, local and remote, through CLI and REST



12

Docker

Luca Cabibbo ASW



Registry of images

- A *registry* is a service (public or private) that contains a collection of container images
 - *Docker Hub* (<https://hub.docker.com>) is Docker's public registry - but private registries are also possible
 - a *repository* is a portion of a registry that contains a set of container images - usually they are variants or different versions of the same image
- A public registry typically contains *basic images* - which contain only an OS, but in some cases also basic software - but not application software
 - eg, base images are *ubuntu*, *postgres*, *wurstmeister* / *kafka* And *openjdk*



Functionality and use

- Here are the main features offered by the Docker platform
 - create a container (a container instance) from a container image
 - start, monitor, inspect, stop and destroy containers
 - create and manage container images
 - manage related groups of containers - in which to run multi-container distributed applications



* Docker in practice

- Interaction with a Docker host is done through an interface (CLI or remote, the remote interface is based on a REST API)
 - this API is command based **docker** - with numerous options / commands / operations for managing images and containers (and other Docker objects) and their life cycle
 - the commands **docker image** for image management
 - the commands **docker container** for container management
 - some commonly used commands exist in two versions, one extended and one short
 - e.g., **docker container run** And **docker run**
 - e.g., **docker image ls** And **docker images**



Docker in practice

- Some basic Docker commands
 - **docker image build** (or **docker build**) allows you to build a (custom) image
 - **docker build -t *image-name* context**
 - **docker container create** (or **docker create**) allows you to create a new container from an image
 - **docker create --name =*container-name* *image-name***
 - **docker container start** (or **docker start**) allows you to run a container (already created)
 - **docker start *container-name***
 - **docker container run** (or **docker run**) creates and executes a new container (possibly anonymous), using a single command
 - **docker run [--name =*container-name*] *image-name***



- For more examples and ideas, visit: <https://docs.docker.com/engine/userguide/>



-
- The diagram illustrates a neural network architecture. It consists of several layers of nodes connected by lines. The top layer has 10 nodes, followed by a hidden layer with 10 nodes, and another hidden layer with 10 nodes. The output layer has 10 nodes. The connections between layers are represented by lines and symbols like '#', '##', and '=='. The diagram is labeled 'Figure 1: A diagram of a neural network architecture'.



Image construction

- For building a custom image, Docker uses a type approach *infrastructure-as-code* - based on a special text file named **Dockerfile**
 - the **Dockerfile** contains all the commands to execute to build a custom image
 - the command **docker build -t *image-name* *context*** allows you to automatically build an image (named *image-name*) starting from a context *context*
 - the context can be a local folder - in particular, **.** - or a location on a Git repository
 - the context must contain the **Dockerfile**, along with any other files of interest (e.g., binary files, scripts and templates)



Dockerfile - FROM and ENTRYPOINT

- A **Dockerfile** it is made up of a sequence of instructions

```
# Hello world
FROM busybox: latest
ENTRYPOINT ["echo", "Hello, world!"]
```

 - education **FROM** specify the base image from which to build the custom image (and possibly its version)
 - e.g., **busybox** (is a minimal Linux distribution) or **ubuntu: 18.04**
 - **ENTRYPOINT** [***executable***, ***param1***, ***param2***, ...] is a statement that specifies the executable or command to be executed by the containers that will be created from this image
 - a **Dockerfile** it must begin with an education **FROM** and usually ends with a single statement **ENTRYPOINT**



Creation of the image and container

- Building an image
 - `docker build -t myhello`. - from the folder that contains the Dockerfile seen above
 - create a new picture named `myhello`
- Creating a container
 - `docker create --name = myhello myhello`
 - create a new container named `myhello` starting from the image `myhello`
- Running a container
 - `docker start -i myhello`
 - start the container `myhello` (interactively)
 - in this case, display `Hello, world!` and then it ends

Hello, world!



Dockerfile - CMD

- Education **CMD** allows you to specify arguments for the statement **ENTRYPOINT** - with the remark that these arguments can be overridden when the container starts
- ```
Hello world
FROM bosybox: latest
ENTRYPOINT ["echo"]
CMD ["Hello, world!"]
```
- `docker build -t myhello2`.
  - `docker run myhello2`
- Hello, world!
- `docker run myhello2 Hello, world!`
- Hello World!



## Example: Apache HTTP Server

- In the **Dockerfile** other instructions can also be used
  - e.g., the **Dockerfile** for an Apache HTTP server

# Dockerfile for Apache HTTP Server

FROM ubuntu: 18.04

# Install apache2 package

RUN apt-get update && \  
apt-get install -y apache2

# Other instructions

ENV APACHE\_LOG\_DIR / var / log / apache2

VOLUME / var / www / html

EXPOSE 80

# Launch apache2 server in the foreground

ENTRYPOINT ["/usr/sbin/apache2ctl", "-D", "FOREGROUND"]

- now we explain the new instructions



## The RUN instruction

- Education **RUN** specifies a command to run when building an image
  - eg, to request execution of a command or script during the **provisioning** of the container image - and not during the **execution** of the container
  - a **Dockerfile** it can contain multiple instructions **RUN** - which are performed sequentially
- The main difference between education **ENTRYPOINT** and instructions **RUN** it is the moment of their execution
  - the instructions specified by **RUN** are performed during the construction of an image
  - the instruction specified by **ENTRYPOINT** it will be executed by the containers created from the image



## The RUN instruction

- It is usually preferable to have in one **Dockerfile** one instruction **RUN** (or a few) - which specify a sequence of commands separated by **&& \** - instead of lots of instructions **RUN**

- eg

```
Install apache2 package (better!) RUN
apt-get update && \
 apt-get install -y apache2
```

- should be preferred to

```
Install apache2 package (worst!) RUN
apt-get update
RUN apt-get install -y apache2
```

- the explanation for this advice is given below



## Other instructions

- Other instructions for the **Dockerfile**
  - education **COPY *src dest*** copy a set of files or folders from the source ***src*** (which must be relative to the context of the construction of the image) to the destination ***dest*** (in the container)
  - education **ADD *src dest*** it is similar - but allows you to copy remote files (i.e. external to context) into the container
  - education **ENV *key value*** set an environment variable in the container



## The VOLUME statement

- Other instructions for the **Dockerfile**
  - education **VOLUME** *path* defines an external mount point - for data on the host system or in another container
  - education **VOLUME** must be used in conjunction with other command options **docker create** And **docker run**
    - the option **-v** *host-src: container-dest* to mount a host system folder in the container - it is a shared folder between the host and the container
    - the option **-volumes-from** *=container-name* to mount a volume managed by another container into the container



## The EXPOSE statement

- Other instructions for the **Dockerfile**
  - education **EXPOSE** *port* specifies that the container listens at runtime on the port *port*
    - this statement is usually used in conjunction with other command options **docker create** And **docker run** to publish (this is Docker's term for port forwarding) some ports of a container on its host
      - the option **-p** *host-port: container-port* to publish a specific port exposed by the container to a specific port on the host
      - the option **-P.** to publish all ports exposed by the container to random host ports
  - it should be noted that containers can still communicate with each other even on ports that are not exposed or not published on the host



## Example: Apache HTTP Server

### - Dockerfile for an Apache HTTP server

```
Dockerfile for Apache HTTP Server

FROM ubuntu: 18.04

Install apache2 package
RUN apt-get update && \
 apt-get install -y apache2

Other instructions
ENV APACHE_LOG_DIR / var / log / apache2
VOLUME / var / www / html
EXPOSE 80

Launch apache2 server in the foreground
ENTRYPOINT ["/usr/sbin/apache2ctl", "-D", "FOREGROUND"]
```



## Example: Apache HTTP Server

### - Image construction

- `docker build -t myapache.` - from the folder that contains the Dockerfile

### - Container creation

- `docker create`
  - `-v ~/projects/www:/var/www/html -p 8080:80`
  - `- name = myapache myapache`
- the pages served by the HTTP server are those in the host folder `~/projects/www`
- HTTP server is redirected to host port 8080

### - Container execution

- `docker start myapache` - start the container `myapache`
- then you can access the HTTP server from the host on `http://localhost:8080`



## Other Docker commands

- Other useful Docker commands
  - to list containers running (or even stopped)
    - `docker container ls` - or `docker ps [-a]`
  - to inspect the ports used by a container - especially useful when using the option `-P`.
    - `docker container port container-name` - or `docker port`
    - the result is of form `80 / tcp -> 0.0.0.0:8080`
  - to inspect a container or image
    - `docker container inspect container-name` - or `docker inspect`
    - returns container or image information (in JSON format) - e.g., network configuration (including port publishing) and volume sharing



## Other Docker commands

- Other useful Docker commands
  - to view the logs generated in a container
    - `docker container logs container-name` - or `docker logs`
  - to stop a running container
    - `docker container stop container-name` - or `docker stop`
  - to remove a container
    - `docker container rm container-name` - or `docker rm`
  - to stop all running containers (use with caution!)
    - `docker stop $(docker ps -a -q)`
  - to remove all containers (use with caution!)
    - `docker rm $(docker ps -a -q)`



## Other Docker commands

- Other useful Docker commands
  - to list images in the local cache
    - `docker image ls` - or `docker images`
  - to remove an image from the local cache
    - `docker image rm image-name` - or `docker rmi`
  - to remove all images from the local cache (use with caution!)
    - `docker rmi -f $(docker images -q)`



## Other Docker commands

- Other useful Docker commands
  - the Docker client can also be used to specify commands to run on a remote Docker host *docker-host*
    - `docker -H = tcp: //docker-host: 2375 command`
    - `docker -H = tcp: //docker-host: 2376 command`
    - port 2376, unlike 2375, supports secure access over TLS
  - the Docker host must be enabled for remote access
  - alternatively, you can specify the remote Docker host using the environment variable `DOCKER_HOST`
    - `export DOCKER_HOST = tcp: //docker-host: 2375`
    - `docker command` - the command is executed on *dockerhost* instead of locally






## \* How Docker works

- We still discuss how Docker works - specifically, the following aspects
  - format of images and containers
  - construction of images
  - creation of containers
  - running container
  - data sharing (volumes)
  - networks
  - registry



## - Format of images (and containers)

- A key element of Docker is the format used for the image and container file system
  - the file system of an image (or container) is made up of a sequence of layers - each layer is a set of files




- these layers are combined into a single consistent file system using one *Union File System (UFS)*
  - a file is read in the topmost layer in which it is located
  - in a container, the only layer that can be written at runtime is the topmost layer



## Format of images (and containers)

- A key element of Docker is the format used for the image and container file system
  - the file system of an image (or container) is made up of a sequence of layers - each layer is a set of files



## Format of images (and containers)

- In the file system of each image (or container), the base is always a base image - usually it contains an OS and its libraries
  - each subsequent layer usually corresponds to the installation of a package, middleware or application
  - in addition to these layers, each container (but not the images) has a final additional layer, which is the only editable part of the container's file system
    - all writes, modifications and deletions performed in the container operate on the latter additional layer
  - this "light" format
    - allows you to share layers between images and between containers
    - makes it easy to update images (e.g., to update an application to a new version) - which can be done by updating or adding layers, rather than completely rebuilding the images




## Images and containers

- An image



Image

- A container (or rather, its file system)

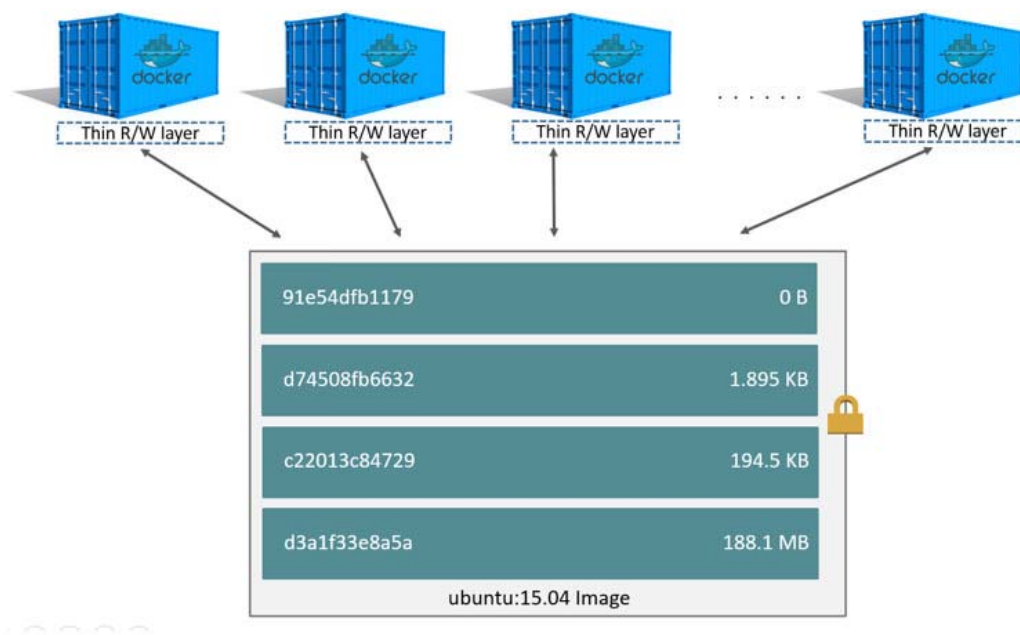


Container  
(based on ubuntu:15.04 image)



## Images and containers

- An image shared by multiple containers





## - Image building

- Building a custom image is based on running a **Dockerfile**, and happens as follows
  - first, the base image specified by **FROM** is downloaded from the registry into an image cache of the host (if it is not already present in the cache)
    - this base image (made up of one or more layers) is used as the base layer (s) of the new custom image
  - then, repeatedly, each instruction of the **Dockerfile** (particularly, **RUN**) is done using a new writable layer on top of the current image
    - the result of executing a del statement **Dockerfile** is then saved (before executing the next instruction)
    - Docker recommends minimizing the number of layers in images and containers - and thus minimizing the number of instructions **RUN** of a **Dockerfile**

41

Docker

Luca Cabibbo ASW



## - Creation of containers

- The creation of a container always takes place from an image
  - a container consists of a file system and meta-data
    - the (layered) file system of the container is obtained from the initial image, to which a new writable layer, specific for the container is added on top - this layer is allocated in the host's file system
  - the images are instead immutable and can be shared by multiple containers

42

Docker

Luca Cabibbo ASW



## - Running containers

### - Running a container

- when a container is requested to run, the container engine allocates runtime resources for the container
  - eg, allocate (in the host kernel) a set of namespaces and configure the network for the container
- then start the container, starting from its file system
- finally, the container executes the command specified by **ENTRYPOINT** (with the arguments specified by **CMD** or from the command line)



## - Volumes and data sharing

- Container storage itself is ephemeral - when a container is destroyed, all of its data is lost
  - how can persistent data be managed?
- A **volume** is a directory outside the container UFS
  - a volume can be accessed, shared and reused by multiple containers
  - a volume allows you to manage persistent data, independently of the life cycle of the individual containers that can access it



## Volumes and data sharing

- A first possibility is to mount a host folder as a volume in a container running on the host, using the option
  - `v` from `docker create` you have `docker run`
    - `docker create -v ~ / projects / www: / var / www / html...`
    - in this case, the data resides on the host (i.e., in a predefined absolute location of the host's file system) - and not in the container
      - therefore, changes to this data are made on the host, persistently
      - an example of use is to redirect the log files of a server running in a container to the host



## Volumes and data sharing

- Another possibility is to have volumes shared between containers - but not tied to a specific folder on the host (in an absolute default location of its file system) - in this case the option must also be used `-volumes-from`
  - a container must first be created, using the option `-v` to indicate a shared folder of the container - but without binding this volume to any folder on the host
    - the volume will reside in this container
    - in practice, this volume resides on the host, but not in an absolute default location
  - then you can create other containers that access that shared volume, with the option `-volumes-from container-name`
  - if the container in which a volume resides is deleted, the volume is still kept (unless an explicit deletion is requested)



- Docker allows you to manage network communication between containers as well as with the host
  - during installation, Docker automatically creates three networks, **bridge**, **host** And it is not - but it is also possible to create others
  - the network **bridge** (in "bridge" mode) is associated with the virtual interface **docker0** on the host and a private network **172.17.0.1/16**
  - when a container is run, Docker associates it with a free IP address on the network **bridge**
    - it is possible to connect a container to a different network using the option **-network =network**
  - containers can communicate with each other knowing the absolute position (IP address and port) of the various services present on the network
  - the network **host** instead, it adds a container to the host's network



- Learn more about networks
  - **docker inspect** allows you to find the information you need to communicate with a container over the network
    - eg, the Apache HTTP server may be exposed to the address **172.17.0.2:80** (of the private network)
  - it is also possible to make these services accessible to the host and outside the host via port mapping (port forwarding) - via options **-p** And **-P**. from **docker create** And **docker run**
    - Docker manages these options by automatically configuring the NAT rules on the host **iptables**
    - the option **-ip** it also allows you to associate a specific IP address (valid for the host) to a container



- Learn more about networks
  - using a *user-defined network* (instead of the network bridge) containers can communicate with each other also through their own "logical" name - as well as by their IP address
    - the container engine operates from DNS for its containers
  - creation of a user-defined network
    - `docker network create -d network-driver network-name`
    - e.g., `docker network create -d bridge my-net`
  - connecting a container to a network
    - `docker run --network = my-net --name = container1 -it busybox`
    - other containers connected to this network can see this container by its "logical" name `container1`
    - a container can also be connected to multiple networks



## - Registry

- A registry is a service for managing a set of container images
  - main operations of a registry
    - `docker pull image-name` - download an image from the registry to the host's local cache - otherwise, `docker build` it does this automatically
    - `docker push image-name` - upload an image to the registry
  - registry query
- Docker's public registry is *Docker Hub* - some of the images it manages are "official"
  - alternatively, *Docker Registry* is a tool for managing your own private registry
  - in the spirit of Docker, Docker Registry can run as a container





## Using Docker Hub

- Using Docker Hub - you must first create your own account, e.g., **aswroma3**
  - login
    - `docker login [-u aswroma3] [-p password] [ server ]`
  - creation and "tagging" of an image
    - `docker build -t aswroma3 / myhello.` or
    - `docker build -t myhello.` followed by `docker tag myhello aswroma3 / myhello`
  - saving an image on the registry (it must be "tagged")
    - `docker push aswroma3 / myhello`
  - loading an image from the registry (optional)
    - `docker pull aswroma3 / myhello`
  - creating and running a container from the image
    - `docker run aswroma3 / myhello`



## Docker Registry



- Managing a (private) Docker Registry - in the spirit of Docker, it can be run as a container
  - start the registry (the `-d` runs the container in the background) - suppose on the node **myregistry**
    - `docker run -d -p 5000: 5000 --restart = always --name registry`
      - `v / var / local / docker / registry: / var / lib / registry registry: 2`
  - creation and "tagging" of an image
    - `docker build -t myhello.`
    - `docker tag myhello myregistry: 5000 / myhello`
  - saving an image on the registry (it must be "tagged")
    - `docker push myregistry: 5000 / myhello`
  - loading an image from the registry
    - `docker pull myregistry: 5000 / myhello`
  - creating and running a container from the image
    - `docker run myregistry: 5000 / myhello`



## - General recommendations

- Some recommendations on containers - and related images
  - only one process per container
    - supports the reuse of images and containers
    - supports horizontal scaling
  - "ephemeral" containers (*ephemeral*, i.e. temporary, passenger, and stateless) - as far as possible
    - so that a container can be stopped and destroyed and then replaced by another container as quickly as possible
  - supports availability and scalability
- minimal containers
  - use the smallest possible base image, avoid installing unnecessary packages, and minimize the number of layers
  - supports availability

53

Docker

Luca Cabibbo ASW



## \* A content-based application

- Before concluding, here is an example of running a simple web application (Spring Boot) in a Docker container
  - the application **lucky-word** - see the handout on Spring Boot
- here is the **Dockerfile** (in the root folder of the Spring Boot project) which uses an image with Open JDK

# Dockerfile for the lucky-word application

FROM openjdk: 11-jdk

# Install the application binary

ADD build / libs / lucky-word.jar lucky-word.jar

EXPOSE 8080

# Launch the Java application

ENTRYPOINT ["java", "-Xmx128m", "-Xms128m"]

CMD ["-jar", "lucky-word.jar"]

54

Docker

Luca Cabibbo ASW



## A containerized application

- Here's how to build and run this application
  - first of all, you need to build the Spring Boot application **lucky-word** (in the development environment)

gradle build

- after that, we need to build a container image for the application (in the environment for Docker)

```
build container image docker
build --rm -t lucky-word.
```

- finally, you have to create and start the container (always in the environment for Docker)

```
run application with default docker run -p 8080:
8080 lucky-word profile
```

```
or, to run the application with the Italian profile docker run -p
8080: 8080 lucky-word
- jar -Dspring.profiles.active = italian lucky-word.jar
```

55

Docker

Luca Cabibbo ASW



## A common mistake

- Be careful to avoid the following common mistakes (which everyone makes sooner or later)
  - after modifying (the source code of) an application, remember (always!) to do the following
    - build (or repeat) the (Java) build of the application
    - build (or repeat) the (Docker) build of the Docker image
    - push (or repeat) the push to Docker Hub of the Docker image (if necessary)
  - Sometimes it may also be necessary to remove the older version of the Docker image from the local cache

56

Docker

Luca Cabibbo ASW



## \* Discussion

- The Docker platform has quickly established itself as the reference technology for containers
  - many companies (including large companies like Google!) use Docker not only for development and testing, but also as a production environment for applications with critical requirements for availability, scalability and elasticity
  - Docker is supported both on premises and in the cloud
  - thanks to Docker, containers have become an alternative and complementary application delivery technology to virtualization
- the rationale for using Docker will be more evident after discussing the composition and orchestration of Docker containers - which is the topic of subsequent lecture notes
- the advice of Sam Newman (author of **Building Microservices**) And
  - "I strongly suggest you give Docker a look "