



- Goals and topics

- Goals

- introduce containers and container-based virtualization
- compare containers and virtual machines
- discuss containers as an option for software release

- Subjects

- introduction
- recalls of preliminary notions
- container-based virtualization
- container
- techniques for container-based virtualization
- introduction to Docker containers
- containers and virtual machines compared
- container and software release
- discussion

3

Container and container-based virtualization

Luca Cabibbo ASW



* Introduction

- Two ways to introduce containers and container-based virtualization

- as a lightweight variant of VMs and system virtualization
- in reference to the use that can be made of it

- Let us consider both points of view

- the most relevant to the software architecture is the use of containers
- we also do a comparison between VMs and containers, and discuss the use of containers in the context of the software release

4

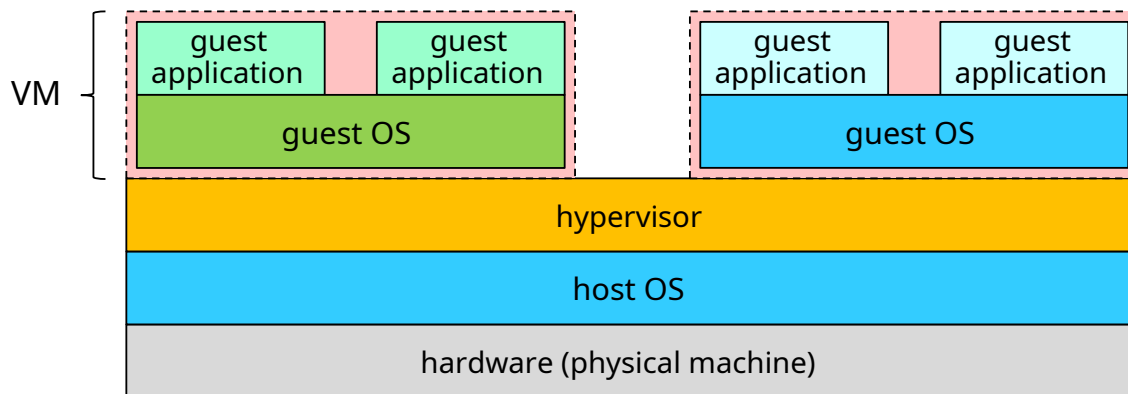
Container and container-based virtualization

Luca Cabibbo ASW



* Recalls of preliminary notions

- System virtualization, based on a hypervisor, offers virtual machine abstraction
 - a VM is a virtual computer - what is virtualized is the hardware of a computer
 - in each VM it is then possible to install a complete OS and run applications and services



5

Container and container-based virtualization

Luca Cabibbo ASW



Preliminary notions

- System virtualization
 - offers several benefits
 - has numerous applications, provides operational flexibility, supports isolation between VMs, each with their own services and applications
 - however, it can introduce a high overhead
 - e.g., in I / O management
 - furthermore, a host hosting N virtual machines must take care of the management and execution of N (instances of) operating systems

6


Container and container-based virtualization

Luca Cabibbo ASW



Preliminary notions

- A **operating system** (*OS*) is made up of several software elements
 - the **kernel** - which handles some critical responsibilities of the OS eg. scheduling, memory manag.
 - a set of **libraries** And **tools** and other utility programs - which operate on top of the kernel es GUI



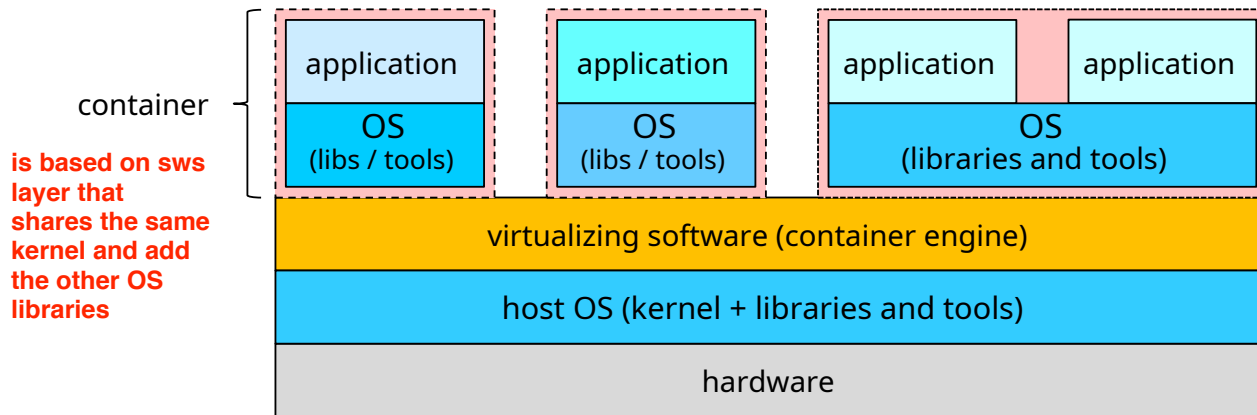
=> idea: create a virtual environment but sharing the host kernel

* Container-based virtualization

- There **virtualization based on container** (*container-based virtualization*)
 - provides the abstraction of **container** ("containers") - also called *lightweight container*
 - a container is a "virtual entity" that gives an abstraction of: comprises the hardware of a computer together with the kernel of an OS
 - in practice, the kernel of each container corresponds to the kernel of the host OS
 - in each container it is then possible to install an OS (libraries and tools) and run applications and services



Container-based virtualization



Implementation

- Some considerations on container-based virtualization
 - it is widespread especially in the Unix / Linux world
 - is a form of lightweight virtualization, also called *OSlevel virtualization* - because it is supported directly by the host OS kernel
 - it does not require a hypervisor on the host
 - it does not use hardware emulation techniques and does not require hardware virtualization support
 - container virtualization software (*container engine*) allows you to define a container as a collection of processes and other resources
 - e.g. a container running a Java service essentially consists of a JVM process (running on the host)



Implementation

- Some considerations on container-based virtualization
 - Only one shared kernel runs on the host system
 - this kernel manages both the resources of the host system and those of the running containers
 - the kernel is shared by the containers
 - but each container runs its own OS and has its own "virtual" resources - e.g., its own complete file system and its own network stack, with its own IP address

can i have containers on multi OS environment? NO eg. i have a hw and an Windows OS
=> i can not make a container linux based (because Windows kernel is not compatible)

But i can make a linux container on a MAC OS machine

Can i make a container that has his own GUI? NO cause the GUI has access on physical layers => containers don't virtualize the hw (VM's do)



Implementation

- Some considerations on container-based virtualization
 - the shared kernel manages all processes and all other resources (eg files), both of the host and of the containers
 - processes of a container are managed as processes of the host
 - can i have a multiprocess program in a container? NO, only in the host
the solution is to program a multithread program or create a VM and than multiprocess
 - a container's file system is managed as a subtree of the host's file system
 - eg, starting with
`/ var / lib / docker / containers / containerid /` filesystem /
 - accessing a file in a container is realized as accessing a file in the host's file system



Implementation

oss: you keep the advantages of virtualizations

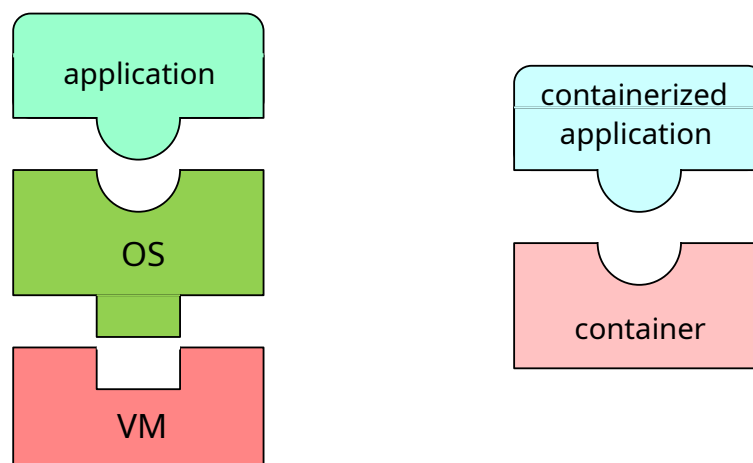
- Some considerations on container-based virtualization
 - container virtualization software ensures isolation between different containers - and between containers and the host
 - e.g., a container cannot interfere with another container's processes or access its file system - nor can it directly access the host's resources
 - however, the containers may not be completely isolated from the host

-flexibility on os versions and libraries



Container and interface


- In terms of interface, comparing VMs and containers
 - the interface exposed by a VM is that of the hardware of a computer
 - the interface exposed by a container is that of an OS's kernel - the system call interface of the kernel





Containers and instances

- A *VM instance* it also includes the OS and the applications that are installed on it
 - similarly, a '*container instance*' it also includes the libraries and tools of the OS and the applications that are installed on it



- the term container is also often used to mean a container instance



Containers and images

- A *VM image* includes the contents of a VM's disks - to facilitate the creation of one or more VM instances from that image
 - similarly, a '*container image*' consists of the file system image of a container
 - includes one or more applications to run in the container
 - along with the libraries and tools of an OS and all the software needed to run those applications
 - so that you can easily create one or more container instances from that image



Discussion

- A preliminary comparison between container and VM
 - containers introduce less overhead than VMs (they are "lighter")
 - performance is almost native - you don't need to use processor virtualization, I / O virtualization, or a hypervisor
 - However, containers offer less operational flexibility and isolation than VMs
 - a container's OS must be compatible with the kernel running on the host



* Container

- Let's now consider containers from the point of view of their use
 - good practice: a container for each application**
 - Each container (container instance) is typically used to run a specific software service or application (although sometimes more than one / one)
 - the container is therefore used to create the virtual execution environment required by that specific software service
 - the container "contains" everything needed to run this software service (which is "contained")



Container

from the developer side:

- A **container** is a standardized software unit, which packages one or more software applications, along with their configurations and dependencies - such that these applications can run quickly and reliably in a suitable container execution environment



Container

- A container is a standardized unit of software, which packages one or more software applications, along with their configurations and dependencies
 - the container has the purpose of providing the application software of interest with a complete and autonomous execution environment, with all the necessary dependencies - without any dependence on the host
 - these dependencies include the OS libraries and tools, the runtime libraries required by the languages used, and the middleware
 - the specific dependencies of the application software of interest are installed and configured in the container (instead of the host)



Container

- A container is a standardized unit of software, which packages one or more software applications, along with their configurations and dependencies
 - the container is also a standardized execution environment for its application (or applications)
 - each application (in its own container) can thus be released and run consistently across a variety of platforms, both on the cloud and on premises
 - There are several standard container formats - the most popular today is Docker



Types of containers

- A classification of containers, in relation to their use
 - a *OS container* is a container designed to be used as a lightweight VM - with its own OS, in which to run multiple applications or services
 - a *application container* is a container designed to hold and run a single application or service
 - we mainly focus on application containers



Application container

- Application containers allow you to focus each container on a single application or service
 - the container must contain only the dependencies for that specific software service
 - this reduces the risk of inconsistencies in the software stack
 - individual containers are lighter - they can be created and started faster, and at runtime they only use the resources required for their specific service
 - this supports availability, scalability and modifiability
 - the containers are isolated from each other
 - this supports reliability and safety
 - from the point of view of an application or service running in a container, it is as if the application or service were running on its own node - with its own IP address and its own file system



* Techniques for container-based virtualization

- There are several technologies for containers, in the context of UNIX and Linux operating systems
 - e.g., LXC, OpenVZ for Linux, Solaris Containers for Solaris, FreeBSD jail for FreeBSD and Docker
 - we describe some container virtualization techniques




- Linux container support

- Some insights into containers *LXC* (*Linux Containers*, 2008) - the first full container implementation for Linux
 - in the Linux kernel, each process can spawn other processes, in a hierarchical fashion
 - a container is a subtree of the system's process tree - associated with resources (such as CPU, memory and disk) and which is kept isolated from other containers (with their resources)
 - a container's kernel is the host's - but a container's OS can be different from the host's OS



LXC container





LXC container

- Practically, *LXC* (*Linux Containers*) is a set of simple tools - but based on a powerful API - to create and manage containers on a Linux host (real or virtual)
 - LXC allows you to create and run one or more containers - with their applications
 - the containers are isolated from each other and from the host OS, and behave as independent machines
 - containers are similar in functionality to VMs
 - but they are controlled directly by the host OS kernel, without the need for a hypervisor



LXC container

- LXC combines some features of the Linux kernel to provide container abstraction - in particular
 - control groups - to control the use of resources
 - namespaces - to control the visibility of resources



Control group

- A *control group* (*cgroup*) And
 - a group of processes - rooted in a certain process and includes all its current and future children (and descendants)
 - associated with parameters and / or limits in the use of resources (such as CPU, memory, network, file system, ...)
 - cgroups allow you to isolate, limit and measure the use of resources assigned to a group of processes



Namespace

- A *namespace* represents a self-contained collection of resources that are given virtual names, which are then mapped onto real resources
 - resource examples are process and user ids, network resources, host name and its ports, files in the file system
 - namespaces allow you to decouple a group of processes from the real resources that will be assigned to it - to control the visibility of resources and to avoid name conflicts and inconsistencies in references
 - in practice, namespaces allow you to separate the resources of different process groups



Discussion

- An LXC container offers an execution environment similar to a standard Linux distribution, with some level of control and resource isolation
 - they are typically used as "OS containers" - to run multiple applications or services as well
- The use of LXC takes place through container management tools
 - which are based on a powerful API - but which is not easy to use




* Introduction to Docker containers

- **Docker** is a container platform
 - to build, release and run distributed applications - simply, quickly, scalable and portable



each app has its Bin/libs OS libraries





Docker container

- A *Docker container* is a standardized unit of software, which packages a software service, along with its configurations and dependencies
 - a container contains everything needed to run that software service - executable code, configurations, libraries, and system tools
 - Docker containers are lightweight, standardized and open and secure



Functionality and use

- Here are the main features offered by the Docker platform
 - create a container (a container instance) from a container image
 - start, monitor, inspect, stop and destroy containers
 - create and manage container images
 - manage related groups of containers - in which to run multi-container distributed applications



Discussion

- Docker (since 2013) has been an instant hit and is used in production by many companies - few technologies have seen a similar adoption rate
 - the main benefits of Docker are lightness, efficiency, simplicity, provisioning speed, openness, possibility of release on a variety of platforms
 - one of the main features of Docker is portability
 - **create a container for each application -> microservices**
 - Docker containers are optimized for the release of individual applications or services - they are "application containers"
 - the ecosystem of tools for Docker is very interesting
 - in particular, it supports the composition and orchestration of containers (discussed in subsequent lecture notes)
 - see also the handout on Docker



* Containers and virtual machines compared

- Containers and virtual machines have characteristics that are complementary to each other
 - VMs are very flexible
 - each VM has its own complete OS and its own applications
 - containers offer less flexibility
 - a container's OS must be compatible with the host's OS (which is usually Unix or Linux)
 - isolation between VMs is complete
 - the isolation offered by containers is not complete
 - however, the flexibility and isolation provided by system virtualization are not always required
 - containers offer execution environments that are adequate for many applications



Containers and virtual machines in comparison

- Containers and virtual machines have characteristics that are complementary to each other
 - Furthermore, the flexibility offered by VMs comes at a cost - one VM
 - requires more resources on the host system
 - it can introduce greater execution overhead
 - it takes longer to start up
 - containers are instead "lighter" than VMs
 - the performance is almost native
 - they require fewer resources
 - eg, a minimal Linux installation takes approx 1.1MB - Ubuntu Server libraries require around 180MB
 - a higher density of containers per host is possible
 - containers can be created and started faster



* Container and software release

- Containers are another release option for distributed software systems
 - each container ("application container") encapsulates a software service, along with the software stack needed for that service



Benefits

- Benefits of using containers for software release
 - each container encapsulates a single software service - releasing an instance of that service can be handled as easily and reliably as creating a container
 - fault isolation and security - each container (with its service) runs quite isolated
 - Containers are lightweight - you can allocate resources to fine-grained containers (and related services)
 - creating and starting a container typically takes from a split second to seconds - less than a VM
 - containers can be released both in the cloud and on premises, in their own private data center
 - specifically, containers can be released into a container orchestration platform (e.g., Kubernetes)



Drawbacks

- Drawbacks to using containers for software release
 - isolation between containers is not complete
 - overhead in administering and updating container images
 - overhead in administering the container execution infrastructure - unless the containers are running in a cloud-hosted solution (such as AWS ECS or Google GCE)



* Discussion

- Both containers and virtual machines have their own advantages and drawbacks
 - VMs offer greater isolation and generality - but at the price of greater overhead
 - containers offer better performance and better resource utilization - but with less isolation and flexibility
 - therefore, containers and VMs have complementary characteristics
 - each technology offers advantages that may be useful in specific situations



Discussion

- It's limiting to think of containers as just a light form of virtualization
 - Containers are significantly changing the way distributed software systems are released and run - and how they are designed and developed
 - the adoption of containers requires a change in software architecture
 - Container adoption is so rapid that regular use of containers in many software systems is expected within a few years
 - "from Gmail to YouTube via Search, all Google products and services run in containers ... every week we run over several billion containers "
 - "80% of all containers in the cloud run on AWS "